

AI Overview

Decision Memo: Missed-Call Recovery, Managed

Verdict: TEST

Score: 67/100 | Confidence: high

Why This Score

The demand is overwhelmingly real (every signal — Avoca AI at a \$1B valuation, \$45–120K/yr per-contractor losses, hot SEO velocity — confirms the pain), and your pricing is well-calibrated to the category. What pulls the score down is structural: this market is now violently saturated, and your only viable launch channel rests entirely on the founder's sales skill. The score is "test, don't build" because the unit economics work *if* one specific assumption holds — and that assumption has not been tested.

Top 3 Strengths

1. Demand (95/100): Pain is quantified by multiple independent sources — 27% of home-services calls go unanswered, average \$1,200–\$5,000 per missed job, \$45K–\$120K/yr per contractor; Beside raised \$20M Nov 2025 specifically for this wedge.
2. Monetization (85/100): Your \$499 setup + \$299/mo + \$15/recovered-job + \$799 concierge is exactly right against Ruby (\$245+/mo) and [Smith.ai](#) (\$292–\$780/mo); LTV ≈ \$3,250, payback < 1 month on cold outreach.
3. Founder-fit on sales motion (58/100): A sales-led founder is the single most useful profile for a B2B managed service — cold pitching, audit delivery, and consultative renewals are exactly the work, and most competing GoHighLevel agencies are run by marketers without sales chops.

Top 3 Risks

1. Competition (32/100): Market saturation is at maximum — 40+ direct AI receptionist vendors, ServiceTitan/Housecall Pro/Jobber bundling native AI receptionist into existing FSM subscriptions, and hundreds of GoHighLevel agencies pitching the same "managed + outcome-based" wedge on TikTok. If ignored: you spend 6 months building a beautiful service that owners can't distinguish from the agency that DM'd them last week, and lose every close-call sale to an incumbent that's already in their stack.
2. Distribution (50/100): You have exactly ONE viable channel (cold outreach + free audit) and zero backup — no audience, no warm intros to trade owners (friends in trades school are not buyers), and \$150/mo budget locks paid acquisition off the table. If ignored: if the cold-to-pilot rate underperforms in your first 30 attempts, you have nothing to fall back on and burn 60–90 days before recognizing it.
3. Retention (58/100): Attribution disputes ("this booking would have come anyway") and the FSM-bundling threat compound on a 12–24 month horizon — clients on Jobber/Housecall Pro will see "free AI receptionist" appear at renewal. If ignored: you build a great month-3 retention story and watch 30–50% of MRR walk at month 12 when their FSM bundle catches up.

Riskiest Assumption

"A sales-led founder with no trades-owner network can convert solo trades to paid pilots via 100% cold outreach + a free 7-day audit at a rate that supports the unit economics — with zero paid acquisition budget."

Test it before building anything. Pick one trade in one metro (plumbing in your home metro is the recommended start — densest pain, highest per-job ACV). Build a 200-shop target list via Google Maps in an afternoon. Cold-contact 50 shops in 14 days via PHONE CALL or in-person at the local supply-house counter (not email — trades ignore it). Pitch the free 7-day audit. Deliver audits using a free-tier Aircall/Google Voice trial + a Notion tracker. Cost: \$0–\$80 plus your 20 hrs/wk. PASS: at least 3 audit acceptances AND at least 1 paid pilot signup from the 50 contacts.

Pre-mortem: If This Fails in 12 Months

1. The cold channel stalls. Audit→pilot conversion comes in below 2% across the first 100 attempts. Three months in, you've delivered 2 pilots and have no paying clients. Side-project willpower runs out before the channel finds its rhythm.
 2. Bundling kills MRR at month 12. You land 8 clients, retention looks great through month 6, then 4 of them hit Jobber Plus renewal and consolidate to the bundled \$39/mo receptionist. ARR halves overnight; LTV assumptions collapse.
 3. Agency competition compresses pricing. Local marketing agencies attached to GoHighLevel start undercutting you with bundled "missed-call + Google Business Profile + reviews" packages at \$399/mo all-in. Your \$299/mo + \$15/job feels expensive and narrow by comparison — close rates fall and you can't differentiate fast enough.
-

What To Do Now

Run the RAT in the next 14 days. Do not write a line of code or commit to tooling beyond free trials. Specifically: pick plumbing in your metro, scrape 200 shops, cold-call or supply-house-walk 50 of them in two weeks, offer the free 7-day call audit, and track the audit-acceptance and pilot-signup numbers. Total spend ≤ \$80. If you hit 3+ audits AND 1+ pilot, you've validated the channel and earned the right to invest in proper tooling and concierge tier setup.

Kill criteria: If you complete 50 contacts in 14 days and get fewer than 3 audit acceptances OR zero paid-pilot signups, the cold-outreach motion does not work for this founder/market combination. Stop and pivot.

If That Doesn't Work

Re-run with one ops change — different trade vertical (HVAC during a seasonal surge) or different motion (in-person supply-house partnership instead of phone outreach). If the second attempt also fails, run [pivot-optimization](#) to explore agency white-label, FSM-consultant referral partnerships, or narrowing to ex-Numa/Goodcall churned customers — these are different distribution models, not just different copy.

What they actually get

1. Free 7-day audit (the door opener)

Before they pay anything, you show them what they're already losing. You track their calls for a week and deliver a one-pager: "You missed 14 calls last week. Based on your average ticket, that's approximately \$8,400 in lost revenue." Their own numbers. Hard to argue with.

2. Never-miss-a-call coverage

- Every call that comes in gets answered — AI handles it if they're on a job, after hours, or just busy

- Caller gets qualified (emergency? type of job? location?), and either gets booked on the spot or gets a text back within minutes
- You handle all the setup, scripting, and ongoing tuning — they touch nothing

3. Human backup layer

When the AI can't handle it (complex job, angry customer, weird request) — a real person steps in. This is the key differentiator over self-serve AI tools that owners configure once and abandon. 67% of contractors who buy AI receptionist tools quit within 90 days because nobody manages the exceptions.

4. Monthly ROI report

Every 30 days they get a one-page report:

- Calls received
- Jobs recovered (AI-handled + human-rescued)
- Estimated dollar value recovered (jobs × their avg ticket)
- Invoice: base + \$15 per recovered job

This report is the reason they stay. Canceling means walking away from a document that says "we made you \$4,200 last month."

The three tiers you pitch

Tier	Price	What they get
Core	\$499 setup + \$299/mo + \$15/job	AI coverage + human escalation + monthly report
Concierge	\$499 setup + \$799/mo	Everything above + live human reply 7am–10pm, priority same-day response
Audit only	Free (7 days)	See what they're losing — no obligation

Always lead with the audit. Nobody says no to free.

What you're NOT selling them

- Software to learn
- A dashboard to log into
- Another subscription they have to manage

- Something that requires them to change how they work

They forward their number. You handle everything. They get a report. That's the entire ask on their end.

The competitive wedge in the pitch

Every plumber has already seen Angi, HomeAdvisor, and three GoHighLevel agency cold calls this month. The difference you lead with:

"Every other service charges you whether they get you jobs or not. We charge \$15 per job we actually recover. If we don't get you jobs, your base fee is \$299 — and the audit will show you exactly what that buys."

Outcome pricing is rare in this space. The research confirmed no major player owns the "pay per recovered job" model — they all charge flat or per-minute. That's your sharpest wedge.

Three things in order — don't skip ahead.

Stage 1 (Weeks 3–6): Prove you can deliver the outcome

Passing the RAT means the channel works. It does NOT yet mean the service works. You converted 1 owner to a pilot. Now you have to actually recover jobs for them and prove the wedge ("zero missed jobs, proven monthly").

- Wire up the real stack: AI receptionist (Numa/Goodcall/Rosie trial) + your script + human-reply backup
- Run the pilot for 30 days, track every missed call → text-back → reply → booked job
- Pass criterion: at end of month, you can hand the owner a one-page report showing ≥5 recovered jobs at ≥\$1,200 avg ticket. That's their ROI proof and your case study.
- Kill criterion: if you can't reliably recover jobs at the rates competitors claim (the 45% text-back success rate is the benchmark), the wedge is fake regardless of how good your sales is.

Stage 2 (Months 2–3): Run the next RAT — can you actually scale the motion?

Your next riskiest assumption (from the ranked list in `scores.json` — criticality 5, uncertainty 4, score 20):

"Solo sales-led founder with no trades network can land 5 pilots in 60–90 days at 20 hrs/wk."

Repeat the playbook from RAT #1 at 4× volume. Goal: 5 paying clients at end of month 3. That's the threshold from your own brief ("land 5 free pilots → case studies → scale"). At $5 \times \$299 = \sim \$1.5K$ MRR. Not life-changing, but enough to prove the funnel exists.

Stage 3 (Months 4–6): Test the moat against GoHighLevel agencies

Next RAT (score 16): does outcome pricing + monthly ROI report actually defend against the agency competition?

- Track churn closely on those first 5 clients
- If anyone churns, do an honest post-mortem (price? attribution dispute? competitor poach?)
- Decision point at month 6: if you have 5+ retained clients at <10% monthly churn → go full-time. If churn is >20% or you can't get past 5 clients → re-run `pivot-engine` (the agency white-label or ex-Numa-churned-customer plays are already pre-identified as alternatives).

Then re-score. After Stage 2, drop your real conversion rate, churn, and unit economics into the system and re-run idea-scoring. The 67 was a pre-launch projection. Real data will move it 15+ points in either direction, and the verdict may upgrade from **test** to **pursue** — or downgrade to **pivot**.